



SQL For Data Science

mm/dd - Steve Mortimer, XXX XXXXXX, XXX XXXXXX

Why Learn SQL?

Fundamental data science skill



- 1. Ubiquitous - It's everywhere in the data community**
- 2. SQL teaches you Relational Algebra and Set Theory**
- 3. Expands Your Computing Power**
- 4. Ease of Access - You can get started today!**

What is a Relational Database?

Key terminology and concepts



- **Rows** represent single items (e.g. an MBA student)
- **Columns** are labeled attributes for each item (e.g. “Name: Steve”, “School: Darden”, “Year: 2”)
- **Tables** are sets of rows that share the same attributes (“Students”)
- **Views** are sets of any rows from different tables in response to a query (“Get me all Second-Year Students in Darden”)
- Proposed in 1970 by E. F. Codd while working at IBM

Why Relational Databases?



Key terminology and concepts

Relational data models:

- Are very easy to extend
- Improve and maintain data integrity
- Allow for re-assembling data many different ways without reorganizing the underlying structure

“Normalization” is the extent to which the data replicated into multiple places within the database.

Highly Normalized = Little Redundancy = More Joins

Denormalized = More Redundancy = Less Joins

Getting Started

The AdventureWorks database



AdventureWorks Cycles is a fictitious bicycle manufacturer

Scenarios include:

- Manufacturing
- Sales
- Purchasing
- Product Management
- Contact Management
- Human Resources

Data needs to be hosted on a running database!



PostgreSQL

AdventureWorks Source:
<https://github.com/infinuendo/AdventureWorks/tree/master/OLTP/Original>

Getting Started

Connecting to the database

Connecting to the database server requires a set of credentials

You will need:

- Host - *ec2-107-22-169-45.compute-1.amazonaws.com*
- Port - *5432*
- Username - *rmfenpdczlrtqo*
- Password - *57a3df647d21f42d4dc029cd951ac7762e1ad440e411dfc8fb7dffb57d82d976*
- Database Name - *d54sjlbs6erhjs*

Check the database is running (open your terminal window):

```
nc -v ec2-107-22-169-45.compute-1.amazonaws.com 5432
```

The diagram shows the command `nc -v ec2-107-22-169-45.compute-1.amazonaws.com 5432` enclosed in a light gray box. Below the command, there are two horizontal brackets. The first bracket spans from the start of the command to the IP address `ec2-107-22-169-45`. The second bracket spans from the port number `5432` to its closing parenthesis. Below the first bracket, the word `HOST` is written in large capital letters. Below the second bracket, the word `PORT` is written in large capital letters.

Getting Started

DBeaver - A SQL client



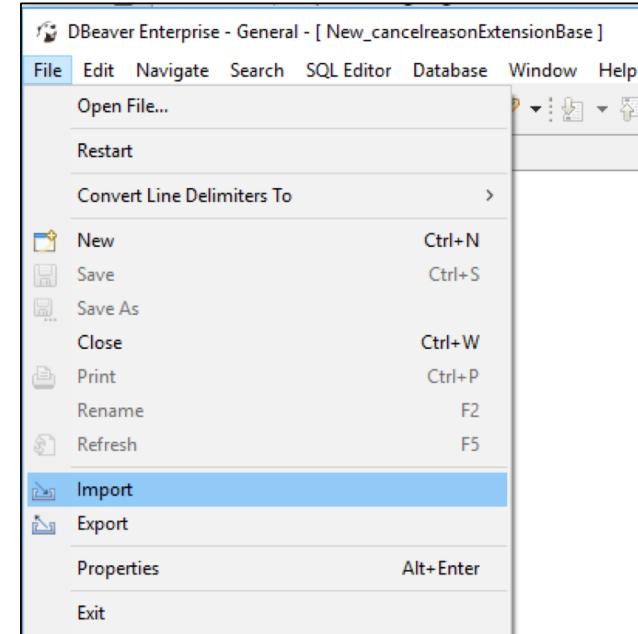
A SQL Client is a program to query databases

First, install DBeaver - <http://dbeaver.jkiss.org/download/>

Second, download [profile](#) (contains connection and SQL scripts)

Third, import the profile into DBeaver

Open up DBeaver -> File -> Import



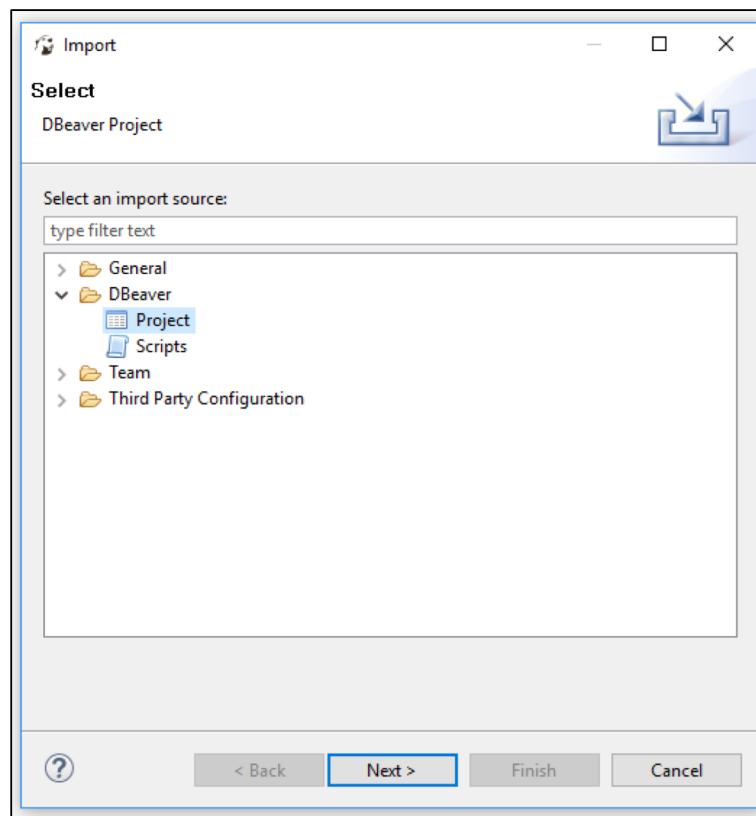
Getting Started

DBeaver - A SQL client (continued)

Next,

Select “Project” as the file type you’d like to import

Click “Next”



Getting Started

DBeaver - A SQL client (continued)

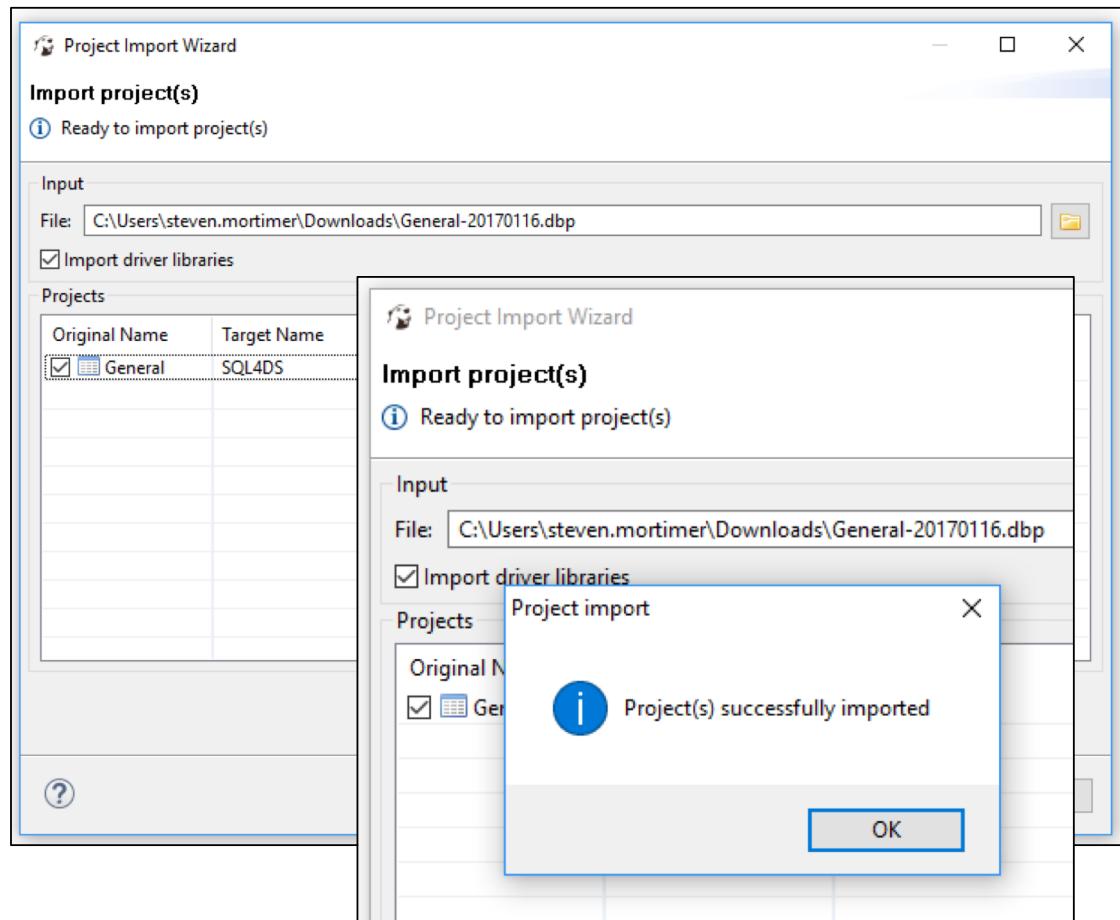
Next,

Select the .dbp file that you downloaded

Ensure that “Import driver libraries” is checked

Ensure the Target Name is “SQL4DS”

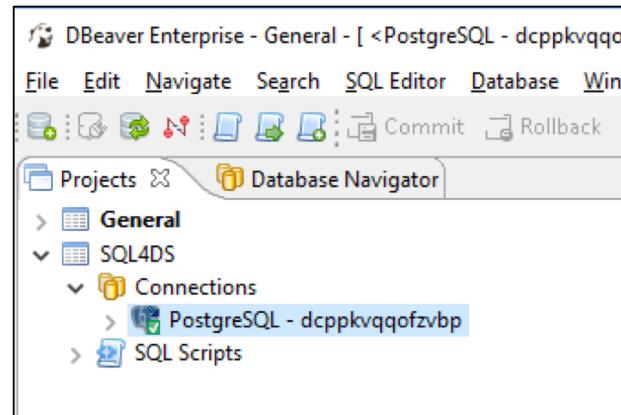
Click “Finish”



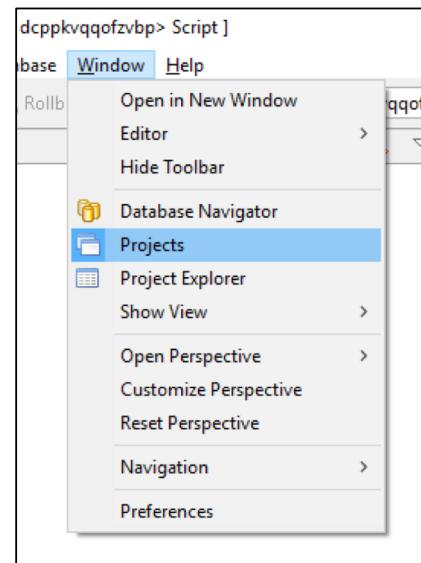
Getting Started

DBeaver - A SQL client (continued)

Click on SQL4DS Project. You should see a connection “PostgreSQL” with green check.



If you do not see a “Projects” tab, go to “Window” -> “Projects”

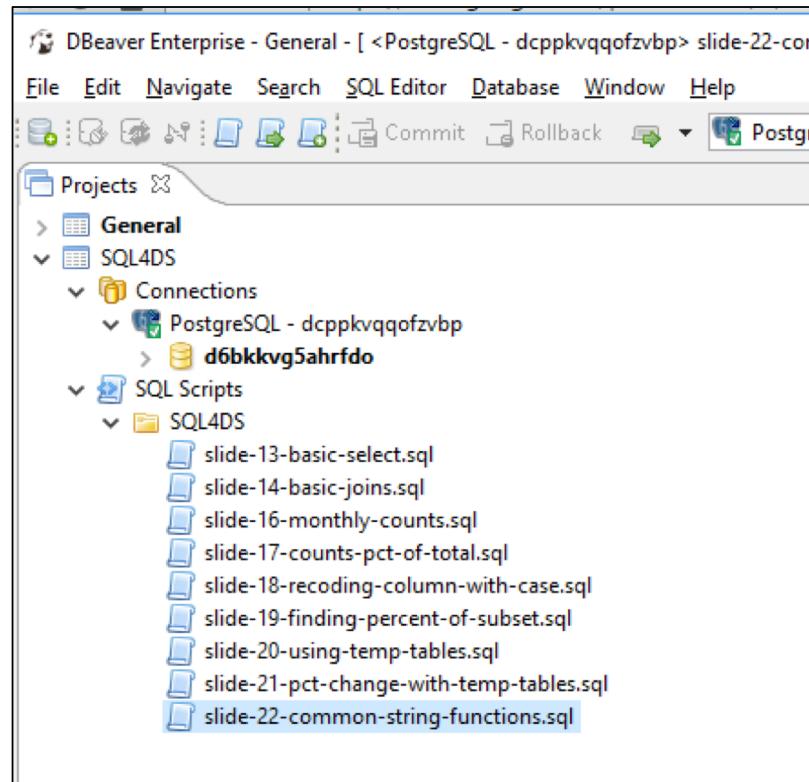


Getting Started

DBeaver - A SQL client (continued)

If you click on arrow to expand the “SQL Scripts” you will see scripts and folders of scripts that are part of your profile.

There is a folder called “SQL4DS” with all scripts for this presentation



Getting Started Checklist



Do you have all the necessary tools?

PostgreSQL (Database)

- Do you have connection credentials?
 - Can you confirm connection via “nc” or “telnet”?
-

DBeaver (SQL Client)

- Installed?
- Do you have project or connection profile?
- Can you connect to database?
- Can you see scripts? Do they execute?

The SQL Basics

SELECT, FROM, WHERE



The basic outline of a SQL query is SELECT, FROM, WHERE

```
SELECT
    productid
    , name
    , listprice
FROM product
WHERE productnumber = 'LO-C100'
;
```

SQL is declarative, just tell the database what you want! But... it can get messy so format your queries to be easy to read.

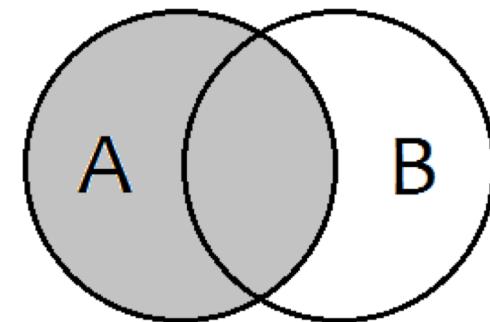
[See the SQL Style Guidelines Cheatsheet](#)

The SQL Basics

JOINS - LEFT JOIN and INNER JOIN

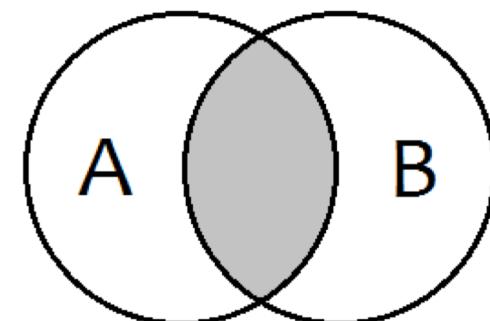
A “JOIN” connects data from 2 or more tables (similar to VLOOKUP)

LEFT JOIN



```
SELECT COUNT(*) -- 504 records
FROM product p
LEFT JOIN production.productcategory as pc
  ON p.productsubcategoryid = pc.productcategoryid
;
```

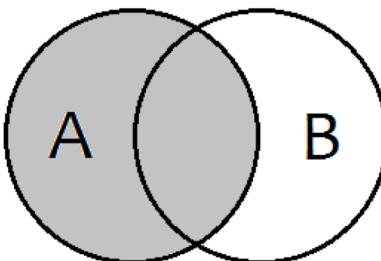
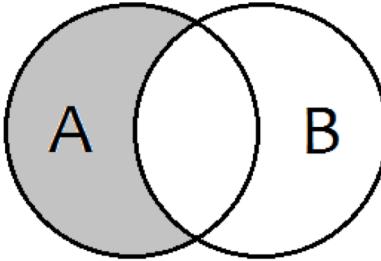
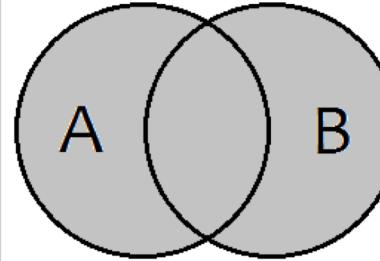
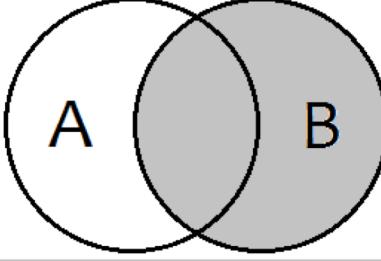
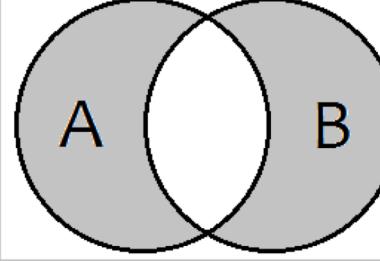
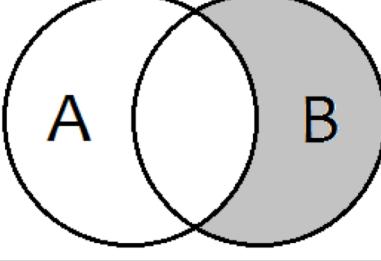
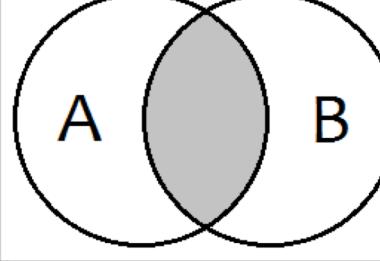
INNER JOIN



```
SELECT COUNT(*) -- 105 records
FROM product p
INNER JOIN production.productcategory as pc
  ON p.productsubcategoryid = pc.productcategoryid
;
```

The SQL Basics

The JOINS cheatsheet

	<pre>SELECT * FROM TableA a LEFT JOIN TableB b ON a.Key = b.Key</pre>	<h2>SQL JOINS</h2>
	<pre>SELECT * FROM TableA a LEFT JOIN TableB b ON a.Key = b.Key WHERE b.Key IS NULL</pre>	 <pre>SELECT * FROM TableA a FULL OUTER JOIN TableB b ON a.Key = b.Key</pre>
	<pre>SELECT * FROM TableA a RIGHT JOIN TableB b ON a.Key = b.Key</pre>	 <pre>SELECT * FROM TableA a FULL OUTER JOIN TableB b ON a.Key = b.Key WHERE a.Key IS NULL OR b.Key IS NULL</pre>
	<pre>SELECT * FROM TableA a RIGHT JOIN TableB b ON a.Key = b.Key WHERE a.Key IS NULL</pre>	 <pre>SELECT * FROM TableA a INNER JOIN TableB b ON a.Key = b.Key</pre>

Find Out Count of Transactions per Month

The GROUP BY statement

```
SELECT    p.name
          , COUNT(*) AS total
FROM transactionhistory th
INNER JOIN product p
  ON th.productid = p.productid
WHERE TO_CHAR(transactiondate, 'YYYY-MM') = '2017-01'
GROUP BY p.name
;
```

Count by month for all months in 2017.

Don't copy/paste the query above! GROUP BY month!

```
SELECT    TO_CHAR(th.transactiondate, 'YYYY-MM') AS month
          , p.name
          , COUNT(*) AS total
FROM transactionhistory th
INNER JOIN product p
  ON th.productid = p.productid
WHERE DATE_PART('year', transactiondate) = '2017'
GROUP BY month, p.name
;
```

Find Out Distribution of Products by Category

Divide by a total count to get the proportion of total

```
SELECT      pc.name AS category
            , COUNT(p.productid)
            , CAST(COUNT(p.productid) AS decimal) / (
                SELECT COUNT(*) FROM production.product) AS pct
FROM product p
INNER JOIN productcategory pc
    ON p.productsubcategoryid = pc.productcategoryid
GROUP BY category
;
```

Be careful! Always inspect your results!!

```
SELECT      pc.name AS category
            , COUNT(p.productid)
            , CAST(COUNT(p.productid) AS decimal) / (
                SELECT COUNT(*) FROM production.product) AS pct
FROM product p
LEFT JOIN productcategory pc
    ON p.productsubcategoryid = pc.productcategoryid
GROUP BY category
;
```

Recoding a Column with CASE

The CASE statement

“CASE” follows *WHEN-THEN-ELSE* logic to change values, here we rename bikes and components to “Core”.

```
SELECT
    p.name
    , CASE
        WHEN pc.name = 'Bikes' OR pc.name = 'Components' THEN
        'Core'
        ELSE 'Non-Core'
    END AS finance_category
FROM product p
LEFT JOIN productcategory pc
    ON p.productsubcategoryid = pc.productcategoryid
;
```

Find Out What Percent of Products are “Core”

Convert a column to 0s and 1s and then the average is the proportion of 1s!

```
SELECT
    AVG(core_indicator) AS pct_core
FROM product p
LEFT JOIN (
    SELECT
        productcategoryid
    , CASE
        WHEN name = 'Bikes' OR name = 'Components' THEN 1
        ELSE 0
    END AS core_indicator
    FROM productcategory
) pc
ON p.productsubcategoryid = pc.productcategoryid
;
```

Note: The product table is LEFT JOIN’ed to a “subquery”. A subquery is its own query that is part of a larger query.

Using Temp Tables

Smaller, separate tables to make large queries more manageable

```
WITH t AS (
    SELECT productid
          , TO_CHAR(transactiondate, 'YYYY-MM') AS month
     FROM transactionhistory
    WHERE DATE_PART('year', transactiondate) = '2017'
)

, t2 AS (
    SELECT t.* , p.name
      FROM t
     INNER JOIN product p
       ON t.productid = p.productid
)

SELECT month, name, COUNT(*) AS total
  FROM t2
 GROUP BY month, name
 ORDER BY month, name
;
```

Find Out the Percent Change Month over Month

The LAG() function

```
, monthly AS (
    SELECT month, COUNT(*) AS total
    FROM t2
    GROUP BY month
    ORDER BY month
)
, lagged AS (
    SELECT
        month
        , total
        , LAG(total) OVER (ORDER BY month) AS prior_month
    FROM monthly
)

SELECT month, total, prior_month,
       (total::float - prior_month) / prior_month AS pct_diff
FROM lagged
;
```

Manipulate Strings and Find via Fuzzy Match



Text columns can be modified and searched as needed

```
SELECT
    name AS original
    , LENGTH(name) AS str_length
    , UPPER(name) AS str_upper
    , SUBSTRING(name, 1, 3) AS str_sub
    , REPLACE(name, 'a', 'zzz') AS str_replace
    , CONCAT('PREFIX_', name) AS str_concat
FROM product
WHERE productnumber LIKE 'L0%'
;
```

LENGTH, UPPER, SUBSTRING, REPLACE, CONCAT will modify a column. LIKE will perform a fuzzy match

Resources

Follow these links for more information on learning SQL

W3Schools - <http://www.w3schools.com/sql/>

Hacker Rank - <https://www.hackerrank.com/domains/sql/select>

SQL Zoo - <http://sqlzoo.net/wiki/AdventureWorks>

Codecademy - <https://www.codecademy.com/learn/learn-sql>

Other Data Science Club Members

Experienced Professionals are all around you!

Questions?



In Person - Ask Now!

Via Email -

Club Inbox - DataScienceClub@darden.virginia.edu

Steve Mortimer - MortimerS19@darden.virginia.edu

**DATA
SCIENCE
CLUB**

